



SUPPORT VECTOR MACHINE

The idea of support vectors and its
importance



Introduction

- The support vector machine is currently considered to be the best off-the-shelf learning algorithm and has been applied successfully in various domains.
- Support vector machines were originally designed for binary classification.
- Then, it is extended to solve multi-class and regression problems.
- But, it is widely used in classification objectives.
- The objective of the support vector machine algorithm is to find a hyperplane in an N -dimensional space (N — the number of features) that distinctly classifies the data points.


- To separate the two classes of data points, there are many possible hyperplanes that could be chosen. Our objective is to find a plane that has the maximum margin, i.e the maximum distance between data points of both classes.



- Maximizing the margin distance provides some reinforcement so that future data points can be classified with more confidence.




- Maximizing the margin distance provides some reinforcement so that future data points can be classified with more confidence.
- The goal is to choose a hyperplane with the greatest possible margin between the hyperplane and any point within the training set, giving a greater chance of new data being classified correctly.




- Hyperplanes are decision boundaries that help classify the data points.
- Data points falling on either side of the hyperplane can be attributed to different classes.
- Also, the dimension of the hyperplane depends upon the number of features.
- It becomes difficult to imagine when the number of features exceeds 3.


A hyperplane in \mathbb{R}^2 is a line




A hyperplane in \mathbb{R}^3 is a plane



- Support vectors are data points that are closer to the hyperplane and influence the position and orientation of the hyperplane.
- Using these support vectors, we maximize the margin of the classifier.
- Deleting the support vectors will change the position of the hyperplane.
- These are the points that help us build our SVM (that works for a new/test data)




- But what happens when there is no clear hyperplane?
- A dataset will often look more like the jumbled balls below which represent a linearly non separable dataset.



- In order to classify a dataset like the one above it's necessary to move away from a 2d view of the data to a 3d view.

- Explaining this is easiest with another simplified example.
- Imagine that our two sets of colored balls above are sitting on a sheet and this sheet is lifted suddenly, launching the balls into the air.
- While the balls are up in the air, you use the sheet to separate them.
- This 'lifting' of the balls represents the mapping of data into a higher dimension.



- This is known as **kernelling**.

Pros & Cons of Support Vector Machines

Pros

- Accuracy
- Works well on smaller cleaner datasets
- It can be more efficient because it uses a subset of training points

Cons

- Isn't suited to larger datasets as the training time with SVMs can be high
- Less effective on noisier datasets with overlapping classes

Applications

- SVM is used for text classification tasks such as category assignment, detecting spam and sentiment analysis.
- It is also commonly used for image recognition challenges, performing particularly well in aspect-based recognition and color-based classification.
- SVM also plays a vital role in many areas of handwritten digit recognition, such as postal automation services.

Derivation of Support Vector Equation

Comparison with logistic regression

Training set: $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$

m examples $x \in \begin{bmatrix} x_0 \\ x_1 \\ \dots \\ x_n \end{bmatrix}$
 $x_0 = 1, y \in \{0, 1\}$


$$\log(\text{odds}) = \log \frac{P(\text{Class 1}|\mathbf{x})}{1 - P(\text{Class 1}|\mathbf{x})} = w_0 + w_1 x_1 + \dots + w_n x_n$$

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

How to choose parameters θ ?

Max Likelihood Estimation (**already discussed**)

Sigmoid function




Threshold classifier output $h_{\theta}(x)$ at 0.5:

If $h_{\theta}(x) \geq 0.5$, predict “y = 1”

If $h_{\theta}(x) < 0.5$, predict “y = 0”

Comparison with logistic regression

- In SVM, we take the output of the linear function and if that output is greater than 1, we identify it with one class and if the output is less than -1, we identify it with another class.
- Since the threshold values are changed to 1 and -1 in SVM, we obtain this reinforcement range of values $[-1,1]$ which acts as margin.



Sigmoid function


Threshold classifier output $h_{\theta}(x)$ at 0.5:

If $h_{\theta}(x) \geq 0.5$, predict “y = 1”

If $h_{\theta}(x) < 0.5$, predict “y = 0”

Comparison with logistic regression

- In SVM, we take the output of the linear function and if that output is greater than 1, we identify it with one class and if the output is less than -1, we identify it with another class.
- Since the threshold values are changed to 1 and -1 in SVM, we obtain this reinforcement range of values $[-1,1]$ which acts as margin.



$$g(\mathbf{x}) = w_1 x_1 + w_2 x_2 + w_0 = 0$$

- $g(x)$ is a linear discriminant function that divides (categorizes) \mathcal{R}^2 into two decision regions.

- The generalization of the linear discriminant function for an n-dimensional feature space in \mathfrak{R}^n is straight forward:

$$g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0 = 0$$

$\mathbf{x} = [x_1 \ x_2 \ \dots \ x_n]^T$ is the feature vector

$\mathbf{w} = [w_1 \ w_2 \ \dots \ w_n]^T$ is a *weight vector*


$w_0 = \textit{bias}$ parameter

- The discriminant function is now a linear n-dimensional surface, called a hyperplane; symbolized as \mathcal{H}
- A two-category classifier implements the following decision rule:

Decide Class 1 if $g(x) > 0$ and Class 2 if $g(x) < 0$


- Thus, \mathbf{x} is assigned to Class 1 if the inner product $\mathbf{w}^T \mathbf{x}$ exceeds the threshold (bias) $-w_0$, and to Class 2 otherwise.

- Figure shows the architecture of a typical implementation of the linear classifier.
- It consists of two computational units: an aggregation unit and an output unit.



A simple linear classifier

- Geometry for $n = 2$ with $w_1 > 0$, $w_2 > 0$ and $w_0 < 0$ is shown in Figure below.
- The origin is on the negative side of \mathcal{H} if $w_0 < 0$, and if $w_0 > 0$, the origin is on the positive side of \mathcal{H} .
- If $w_0 = 0$, the hyperplane passes through the origin.



Linear decision boundary between two classes

Location of any point \mathbf{x} may be considered relative to \mathcal{H} .

Defining \mathbf{x}_p as the normal projection of \mathbf{x} onto \mathcal{H} ,

$$\mathbf{x} = \mathbf{x}_p + r \frac{\mathbf{w}}{\|\mathbf{w}\|}$$

where $\|\mathbf{w}\|$ is the Euclidean norm of \mathbf{w} and $\frac{\mathbf{w}}{\|\mathbf{w}\|}$ is a unit vector.

$$g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0 = \mathbf{w}^T \left(\mathbf{x}_p + r \frac{\mathbf{w}}{\|\mathbf{w}\|} \right) + w_0$$

$$= \mathbf{w}^T \mathbf{x}_p + w_0 + \frac{\mathbf{w}^T r \mathbf{w}}{\|\mathbf{w}\|}$$

$$= r \frac{\mathbf{w}^T \mathbf{w}}{\|\mathbf{w}\|} = r \frac{\|\mathbf{w}\|^2}{\|\mathbf{w}\|} = r \|\mathbf{w}\|$$


**Algebraic measure of the
distance from \mathbf{x} to the hyperplane**

$$r = \frac{g(\mathbf{x})}{\|\mathbf{w}\|}$$

As

$$g(\mathbf{x}) = r \cdot \|\mathbf{w}\|$$

$|g(\mathbf{x})|$ is a measure of the Euclidean distance of the point \mathbf{x} from the decision hyperplane \mathcal{H} .




$$g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0 \begin{cases} > 0 & \text{if } \mathbf{x} \in \mathcal{H}^+ \\ = 0 & \text{if } \mathbf{x} \in \mathcal{H} \\ < 0 & \text{if } \mathbf{x} \in \mathcal{H}^- \end{cases}$$

$$g(\mathbf{x}_d) = \mathbf{w}^T \mathbf{x}_d + w_0 = 0; \mathbf{x}_d = d \frac{\mathbf{w}}{\|\mathbf{w}\|}$$

$$d \frac{\mathbf{w}^T \mathbf{w}}{\|\mathbf{w}\|} + w_0 = 0; d \frac{\|\mathbf{w}\|^2}{\|\mathbf{w}\|} = -w_0; d = \frac{-w_0}{\|\mathbf{w}\|}$$

Perpendicular distance d from coordinate origin to $\mathcal{H} = w_0 / \|\mathbf{w}\|$


Geometry for 3-dimensions ($n=3$)




Hyperplane \mathcal{H} separates the feature space into two half space \mathcal{H}^+ and \mathcal{H}^-

Linear Maximal Margin Classifier for Linearly Separable Data

- For linearly separable, many hyperplanes exist to perform separation.
- SVM framework tells which hyperplane is best.
- Hyperplane with the largest *margin which* minimizes training error.
- Select the decision boundary that is far away from both the classes.
- Large margin separation is expected to yield good generalization.
- in $w^T x + w_0 = 0$, w defines a direction perpendicular to the hyperplane.
- w is called the normal vector (or simply normal) of the hyperplane.
- Without changing the normal vector w , varying w_0 moves the hyperplane parallel to itself.




(a) Large margin separation



(b) Small margin separation

Large margin and small margin separation

Two parallel hyperplanes \mathcal{H}_1 and \mathcal{H}_2 that pass through $\mathbf{x}^{(i)}$ and $\mathbf{x}^{(k)}$ respectively.



Geometric interpretation of algebraic distances of points to a hyperplane for two-dimensional case

\mathcal{H}_1 and \mathcal{H}_2 are parallel to the hyperplane $\mathbf{w}^T \mathbf{x} + w_0 = 0$.

$$\mathcal{H}_1: \mathbf{w}^T \mathbf{x} + w_0 = +1$$

$$\mathcal{H}_2: \mathbf{w}^T \mathbf{x} + w_0 = -1$$

such that

$$\mathbf{w}^T \mathbf{x}^{(i)} + w_0 \geq 1 \quad \text{if } y^{(i)} = +1$$

$$\mathbf{w}^T \mathbf{x}^{(i)} + w_0 \leq -1 \quad \text{if } y^{(i)} = -1$$

or equivalently.

$$d_1 = \frac{1}{\|\mathbf{w}\|}; d_2 = \frac{-1}{\|\mathbf{w}\|} \xrightarrow{(\cdot)} (\mathbf{w}^T \mathbf{x}^{(i)} + w_0) \geq 1$$

distance between the two hyperplanes = margin M

$$M = \frac{2}{\|\mathbf{w}\|}$$

This equation states that maximizing the margin of separation between

KKT Condition

Learning problem in SVM

- Linearly separable training examples,


$$\mathcal{D} = \{(\mathbf{x}^{(1)}, y^{(1)}), (\mathbf{x}^{(2)}, y^{(2)}), \dots, (\mathbf{x}^{(N)}, y^{(N)})\}$$

Problem: Solve the following constrained minimization problem:

$$\begin{aligned} & \text{minimize } f(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w} \\ & \text{subject to } y^{(i)} (\mathbf{w}^T \mathbf{x}^{(i)} + w_0) \geq 1; i = 1, \dots, N \end{aligned}$$

This is the formulation of ***hard-margin*** SVM.

Hard margin svm Vs Soft margin svm



Dual formulation of constrained optimization problem:

Lagrangian is constructed:

$$L(\mathbf{w}, w_0, \lambda) = \frac{1}{2} \mathbf{w}^T \mathbf{w} - \sum_{i=1}^N \lambda_i [y^{(i)} (\mathbf{w}^T \mathbf{x}^{(i)} + w_0) - 1]$$

The Karush-Kuhn-Tucker (**KKT**) conditions are as follows:

i. $\frac{\partial L}{\partial \mathbf{w}} = 0 \Rightarrow \mathbf{w} = \sum_{i=1}^N \lambda_i y^{(i)} \mathbf{x}^{(i)}$

$$\frac{\partial L}{\partial w_0} = 0 \Rightarrow \sum_{i=1}^N \lambda_i y^{(i)} = 0$$

ii. $y^{(i)} (\mathbf{w}^T \mathbf{x}^{(i)} + w_0) - 1 \geq 0; i = 1, \dots, N$

iii. $\lambda_i \geq 0; i = 1, \dots, N$

- \mathbf{w} is computed using condition (i) of KKT conditions

$$\mathbf{w} = \sum_{i=1}^N \lambda_i y^{(i)} \mathbf{x}^{(i)}$$

- and w_0 is computed using condition (iv) of KKT conditions

$$\lambda_i [y^{(i)}(\mathbf{w}^T \mathbf{x}^{(i)} + w_0) - 1] = 0; i = 1, \dots, N$$

- from, condition (iii), it can be said that Very small percentage have $\lambda_i > 0$
 - most among N , vanish with $\lambda_i = 0$.
- $\mathbf{x}^{(i)}$ whose $\lambda_i > 0$ are the *support vectors* and they lie on the margin.

- \mathbf{w} is the weighted sum of these training instances that are selected as the support vectors:

$$\mathbf{w} = \sum_{i \in svindex} \lambda_i y^{(i)} \mathbf{x}^{(i)}$$

- where $svindex$ is the set of indices of support vectors
- All support vectors are used to compute w_0 , and then their average is taken for the final value of w_0

$$w_0 = \frac{1}{|svindex|} \sum_{i \in svindex} (y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)})$$

- where $|svindex|$ is the total number of indices in $svindex$, i.e., total number of support vectors.

The majority of λ_i are 0, for which $y^{(i)}(\mathbf{w}^T \mathbf{x}^{(i)} + w_0) > 1$. These are the $\mathbf{x}^{(i)}$ points that exist more than adequately away from the discriminant, and have zero effect on the hyperplane. The instances that are not support vectors have no information; the same solution will be obtained on removing any subset from them. From this viewpoint, the SVM algorithm can be said to be similar to the k -NN algorithm (Section 3.4) which stores only the instances neighboring the class discriminant.

During testing, we do not enforce a margin. We calculate

$$g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$


and choose the class according to the sign of $g(\mathbf{x})$: $\text{sgn}(g(\mathbf{x}))$ which we call the *indicator function* i_F ,

$$i_F = \hat{y} = \text{sgn}(\mathbf{w}^T \mathbf{x} + w_0)$$

Choose Class 1 ($\hat{y} = +1$) if $\mathbf{w}^T \mathbf{x} + w_0 > 0$, and Class 2 ($\hat{y} = -1$) otherwise.

Linear Soft Margin Classifier for Overlapping Classes

- To generalize SVM, allow noise in the training data.
- Hard margin linear SVM algorithm will not work.



Soft decision boundary

- To allow error in data, relax margin constraints by invoking *slack* variables $\zeta_i (\geq 0)$:

$$\mathbf{w}^T \mathbf{x}^{(i)} + w_0 \geq 1 - \zeta_i \text{ for } y^{(i)} = +1$$

$$\mathbf{w}^T \mathbf{x}^{(i)} + w_0 \leq -1 + \zeta_i \text{ for } y^{(i)} = -1$$

Thus, new constraints:

$$y^{(i)}(\mathbf{w}^T \mathbf{x}^{(i)} + w_0) \geq 1 - \zeta_i; i = 1, \dots, N$$

$$\zeta_i \geq 0$$

Penalize the errors by assigning extra cost and change the objective function to

$$\frac{1}{2} \mathbf{w}^T \mathbf{w} + C \left(\sum_{i=1}^N \zeta_i \right); C \geq 0$$

- Hence it all boils down to optimization problem

$$\begin{aligned} & \text{minimize} \quad \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{i=1}^N \zeta_i \\ & \text{subject to} \quad y^{(i)} (\mathbf{w}^T \mathbf{x}^{(i)} + w_0) \geq 1 - \zeta_i; i = 1, \dots, N \\ & \quad \quad \quad \zeta_i \geq 0; i = 1, \dots, N \end{aligned}$$

This formulation is the *soft margin* SVM.

Lagrangian

$$L(\mathbf{w}, w_0, \zeta, \lambda, \mu) = \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{i=1}^N \zeta_i - \sum_{i=1}^N \lambda_i [y^{(i)} (\mathbf{w}^T \mathbf{x}^{(i)} + w_0) - 1 + \zeta_i] - \sum_{i=1}^N \mu_i \zeta_i$$

Using KKT conditions, the dual formulation of the *soft-margin SVM* is reduced to

$$\text{maximize } L_*(\lambda) = \sum_{i=1}^N \lambda_i - \frac{1}{2} \sum_{i=1}^N \sum_{k=1}^N \lambda_i \lambda_k y^{(i)} y^{(k)} \mathbf{x}^{(i)T} \mathbf{x}^{(k)}$$

$$\text{subject to } \sum_{i=1}^N \lambda_i y^{(i)} = 0$$

$$0 \leq \lambda_i \leq C; i = 1, \dots, N$$

- ζ_i and μ_i - not in the dual objective function.
- The objective function is identical to that for separable case.
- Only difference - constraint $\lambda_i \leq C$

λ_i can have values in the interval $0 \leq \lambda_i \leq C$. thus three cases are there:

Case 1: $\lambda_i = 0$

Don't contribute to the optimum value of \mathbf{w} .

Case 2: $0 < \lambda_i < C$

Corresponding patterns are on the margin.

Case 3: $\lambda_i = C$

Corresponding pattern is misclassified or lies inside the margin.

- support vectors define \mathbf{w} ($\lambda_i > 0$); $\mathbf{w} = \sum_{i \in svindex} \lambda_i y^{(i)} \mathbf{x}^{(i)}$

svindex is the set of indices of support vectors

$$\mathbf{w} = \sum_{i \in svindex} \lambda_i y^{(i)} \mathbf{x}^{(i)}$$


Kernel Function: Dealing with non linearity

Non-linear classifiers

- for several real-life datasets, the decision boundaries are nonlinear.
- To deal with nonlinear case, the formulation and solution methods employed for the linear case are still applicable.
- Only input data is transformed from its original space into another space (higher dimensional space) so that a linear decision boundary can separate Class 1 examples from Class 2.
- The transformed space is called the feature space.
- The original data space is known as the input space.

Non-linear classifiers

- For training examples which cannot be linearly separated.
- In the feature space, they can be separated linearly with some transformations.



Transformation from input space to feature space

The new optimization problem becomes

$$\begin{aligned}
 & \text{minimize} \quad \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{i=1}^N \zeta_i \\
 & \text{subject to} \quad y^{(i)} (\mathbf{w}^T \boldsymbol{\Phi}(\mathbf{x}^{(i)}) + w_0) \geq 1 - \zeta_i; i = 1, \dots, N \\
 & \quad \quad \quad \zeta_i \geq 0; i = 1, \dots, N
 \end{aligned}$$

The corresponding dual is

$$\begin{aligned}
 & \text{minimize} \quad L_*(\lambda) = \sum_{i=1}^N \lambda_i - \frac{1}{2} \sum_{i=1}^N \sum_{k=1}^N \lambda_i \lambda_k y^{(i)} [\boldsymbol{\Phi}(\mathbf{x}^{(i)})]^T \boldsymbol{\Phi}(\mathbf{x}^{(k)}) \\
 & \text{subject to} \quad \sum_{i=1}^N \lambda_i y^{(i)} = 0
 \end{aligned}$$

The decision boundary becomes:

$$\sum_{i=1}^N \lambda_i y^{(i)} [\boldsymbol{\Phi}(\mathbf{x}^{(i)})]^T \boldsymbol{\Phi}(\mathbf{x}) + w_0 = 0$$

- Is there a need to know the mapping of $\boldsymbol{\Phi}$? No.

In SVM, this is done through the use of *kernel function*, denoted by K .

$$K(\mathbf{x}^{(i)}, \mathbf{x}) = [\boldsymbol{\Phi}(\mathbf{x}^{(i)})]^T \boldsymbol{\Phi}(\mathbf{x})$$

There is no explicit need to know what $\boldsymbol{\Phi}$ is.

Constructing Kernels:

- Does any kernel work? No, only valid kernel functions work. Identification of Φ is not needed if it can be shown whether the function is a kernel or not without the need of mapping.
- Function satisfying Mercer's theorem can work as kernel function.
- Mercer's theorem, which provides a test whether a function $K(\mathbf{x}^{(i)}, \mathbf{x}^{(k)})$ constitutes a valid kernel without having to construct the function $\Phi(\mathbf{x})$.

$$K(\mathbf{x}^{(i)}, \mathbf{x}^{(k)}) = [\Phi(\mathbf{x}^{(i)})]^T \Phi(\mathbf{x}^{(k)})$$

Mercer's theorem:

Every semi-positive definite symmetric function is a kernel

$K(\mathbf{x}^{(i)}, \mathbf{x}^{(k)})$ is a kernel function if and only if the matrix \mathbf{K} is positive semidefinite.

$$\mathbf{K} = \begin{bmatrix} K(\mathbf{x}^{(1)}, \mathbf{x}^{(1)}) & K(\mathbf{x}^{(1)}, \mathbf{x}^{(2)}) & \dots & K(\mathbf{x}^{(1)}, \mathbf{x}^{(N)}) \\ \vdots & \vdots & & \vdots \\ K(\mathbf{x}^{(N)}, \mathbf{x}^{(1)}) & K(\mathbf{x}^{(N)}, \mathbf{x}^{(2)}) & \dots & K(\mathbf{x}^{(N)}, \mathbf{x}^{(N)}) \end{bmatrix}$$

A positive semidefinite matrix is a Hermitian matrix (a complex square matrix that is equal to its own conjugate transpose) all of whose eigenvalues are nonnegative.

Polynomial and Radial Basis Kernel

Common kernel functions used:

Polynomial kernel of degree d

$$K(\mathbf{x}^{(i)}, \mathbf{x}^{(k)}) = (\mathbf{x}^{(i)T} \mathbf{x}^{(k)} + c)^d; c > 0, d \geq 2$$


Gaussian radial basis function kernel (RBF)

$$K(\mathbf{x}^{(i)}, \mathbf{x}^{(k)}) = \exp \left(-\frac{\|\mathbf{x}^{(i)} - \mathbf{x}^{(k)}\|^2}{2\sigma^2} \right); \sigma > 0$$

Each of these results in a different nonlinear classifier in (the original) input space.

Polynomial Kernel

Using a polynomial kernel they can be linearly separated in 3D space



Polynomial Kernel

- The polynomial kernel represents the similarity of vectors (training samples) in a feature space over polynomials of the original variables, allowing learning of non-linear models.
- It looks not only at the given features of input samples to determine their similarity, but also combinations of these (interaction features).
- Quite popular in natural language processing (NLP).
- The most common degree is $d = 2$ (quadratic), since larger degrees tend to overfit on NLP problems.
- One problem with the polynomial kernel is that it may suffer from numerical instability: (result ranges from 0 to infinity)

Radial Basis Kernel

- RBF kernels are the most generalized form of kernelization.
- It is one of the most widely used kernels due to its similarity to the Gaussian distribution.
- The RBF kernel function for two points X_1 and X_2 computes the similarity or how close they are to each other.

$$K(\mathbf{x}^{(i)}, \mathbf{x}^{(k)}) = \exp\left(-\frac{\|\mathbf{x}^{(i)} - \mathbf{x}^{(k)}\|^2}{2\sigma^2}\right); \sigma > 0$$

where,


‘ σ ’ is the variance and our hyperparameter

$\|\mathbf{x}_1 - \mathbf{x}_2\|$ is the Euclidean (L_2 -norm) Distance between two points X_1 and X_2

Radial Basis Kernel

- The maximum value that the RBF kernel can get is 1 and occurs when d_{12} is 0 which is when the points are the same, i.e. $X_1 = X_2$.
- When the points are the same, there is no distance between them and therefore they are extremely similar.
- When the points are separated by a large distance, then the kernel value is less than 1 and close to 0 which would mean that the points are dissimilar.

- There are no golden rules for determining which admissible kernel will result in the most accurate SVM.
- In practice, the kernel chosen does not generally make a large difference in resulting accuracy.
- SVM training always finds a global solution, unlike neural networks (to be discussed in the next chapter) where many local minima usually exist.



Kernel ϕ

